


# Communication in Distributed Systems






# Outline

- Interprocess communication
- Client server model
- distributed system, it is completely different from uniprocessor system as there is no shared memory.
- 






# Interprocess Communication

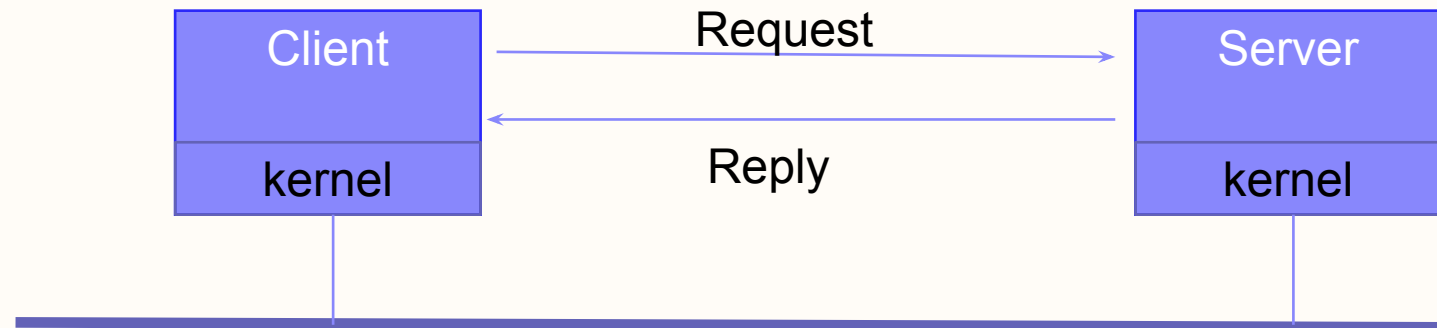
- In distributed system, it is completely different from uniprocessor system as there is no shared memory.
- Certain rules need to be followed for interprocess communication called Protocols.
- For wide area distributed systems, these protocols take the form of multiple layers such as OSI , ATM.
- OSI model addresses only a small aspect of the communication - sending bits from the sender to the receiver, with much overheads.






# Client - Server model

- It is based on simple connectionless request / reply protocol.
- Client sends a request message to the server and the server returns the data requested or an error code indicating the reason of failure.
- It is simple. No connection to be established before use and no connection to be closed after use.
- Simplicity leads to efficiency. Only three levels of protocol are needed.





# Client - Server model

- Physical and datalink protocol take care of getting the packets from client to server and back.
- No routing and no connections - layers 3 & 4 not needed.
- Layer 5 is the request/ reply protocol. No sessions required.
- Communication provided by the micro-kernels using two system calls -
  - *send (dest, &mptr)*
  - *receive(addr, &mptr)*

mptr - message pointer

dest - destination process

addr - source address

7

6


5 Request / reply

4

3

2 Datalink

1 Physical



# Example - client and server


```
/* Definitions needed by clients and servers. */
#define TRUE 1
#define MAX_PATH 255 /* maximum length of file name */
#define BUF_SIZE 1024 /* how much data to transfer at once */
#define FILE_SERVER 243 /* file server's network address */

/* Definitions of the allowed operations */
#define CREATE 1 /* create a new file */
#define READ 2 /* read data from a file and return it */
#define WRITE 3 /* write data to a file */
#define DELETE 4 /* delete an existing file */

/* Error codes. */
#define OK 0 /* operation performed correctly */
#define E_BAD_OPCODE -1 /* unknown operation requested */
#define E_BAD_PARAM -2 /* error in a parameter */
#define E_IO -3 /* disk error or other I/O error */

/* Definition of the message format. */
struct message {
    long source; /* sender's identity */
    long dest; /* receiver's identity */
    long opcode; /* requested operation */
    long count; /* number of bytes to transfer */
    long offset; /* position in file to start I/O */
    long result; /* result of the operation */
    char name[MAX_PATH]; /* name of file being operated on */
    char data[BUF_SIZE]; /* data to be read or written */
};
```

header.h




# Example - client and server

```
#include <header.h>
void main(void) {
    struct message m1, m2;          /* incoming and outgoing messages */
    int r;                          /* result code */

    while(TRUE) {                  /* server runs forever */
        receive(FILE_SERVER, &m1); /* block waiting for a message */
        switch(m1.opcode) {        /* dispatch on type of request */
            case CREATE: r = do_create(&m1, &m2); break;
            case READ:   r = do_read(&m1, &m2); break;
            case WRITE:  r = do_write(&m1, &m2); break;
            case DELETE: r = do_delete(&m1, &m2); break;
            default:     r = E_BAD_OPCODE;
        }
        m2.result = r;              /* return result to client */
        send(m1.source, &m2);      /* send reply */
    }
}
```

A sample server.



# Example - client and server

```
(a)
#include <header.h>
int copy(char *src, char *dst){
    struct message ml;
    long position;
    long client = 110;


    initialize( );
    position = 0;
    do {
        ml.opcode = READ;
        ml.offset = position;
        ml.count = BUF_SIZE;
        strcpy(&ml.name, src);
        send(FILESERVER, &ml);
        receive(client, &ml);

        /* Write the data just received to the destination file.
        ml.opcode = WRITE;
        ml.offset = position;
        ml.count = ml.result;
        strcpy(&ml.name, dst);
        send(FILE_SERVER, &ml);
        receive(client, &ml);
        position += ml.result;
    } while( ml.result > 0 );
    return(ml.result >= 0 ? OK : ml.result);
}
```

/\* procedure to copy file using the server \*/  
/\* message buffer \*/  
/\* current file position \*/  
/\* client's address \*/  
/\* prepare for execution \*/  
/\* operation is a read \*/  
/\* current position in the file \*/  
/\* how many bytes to read\*/  
/\* copy name of file to be read to message \*/  
/\* send the message to the file server \*/  
/\* block waiting for the reply \*/  
/\* operation is a write \*/  
/\* current position in the file \*/  
/\* how many bytes to write \*/  
/\* copy name of file to be written to buf \*/  
/\* send the message to the file server \*/  
/\* block waiting for the reply \*/  
/\* ml.result is number of bytes written \*/  
/\* iterate until done \*/  
/\* return OK or error code \*/

A client using the server to copy a file.






# Addressing


- One way of mentioning server address is to mention it in **header.h** as a constant.
- Sending kernel can extract it (ex - 243 referring to machine) from message structure and use it for sending packets to server.
- Ambiguity arises if multiple processes are running on the same server.

## Alternative 1 -

- Send messages to processes , not machines.
- Process identification - two part names - machine + process no.  
Ex - 243.4 or 4@243
- Each machine can number its processes starting from 0. So there is no confusion between process 'n' of different machines.
- No global coordination is required.

1. Request to 243.0
2. Reply to 199.0





# Addressing




## Alternative 2 -

- use *machine.local-id* instead of *machine.process*
- Each process is assigned a *local-id* and informs kernel that it listens to *local-id*
- Problem - user is aware of the location of the machine (243). If the machine is down, compiled programs with *header.h* will not work, although another machine (365) is available. No transparency.

## Alternative 3-

- Each process has a unique address that doesn't contain machine number.
- A centralized **process address allocator** maintains a counter. Upon receiving a request, it returns the current value of the counter.
- Problem - Such centralized components do not scale to large systems.







# Addressing

## Alternative 4 -

- Each process picks its own identifier from a large address space (space of 64 bit binary integer).
- It is scalable.
- Identification of machine -Sender broadcasts a special **Locate packet** containing the address of the destination process. It will be received by all machines on the network. The matched kernel responds with a message “here i am” alongwith the machine number. So the sending kernel uses this machine number for further communication.
- Problem - Broadcasting is an overload to the system.

1. Broadcast
2. Here I am
3. Request
4. Reply





# Addressing

## Alternative 5 -

- Overload can be avoided by providing an extra machine to map high-level (ASCII) service names to machine address.
- These names are embeded in the programs, not binary machine numbers.
- For the first time client sends a query to the Name server, asking the machine number where the server is currently located. Then the request can be sent directly to the machine address.
- Problem - If the name server is replicated, consistency problem may arise.

1. Lookup
2. NS reply
3. Request
4. Reply

