

# **Replica Management & Consistency Protocols in Distributed Systems**

# 1. What is Replica Management?



- In a distributed system, data is often stored (or replicated) in different places (computers or servers) to make it more reliable, faster, and always available, even if some computers fail.
- Replica management means taking care of these copies (called replicas) to make sure:
- Data is available even if some servers are down.



- The system runs faster by spreading out the workload.
- All copies of the data are the same (consistent).

### 3. What is Consistency?



- When we have multiple copies of data, we want to make sure that all of them show the same information. If you update one copy, the system should update all the others. This is called consistency.

## 4. Different Levels of Consistency




- **Strict Consistency:** All users see the latest data, but it can be slow.
- **Sequential Consistency:** All users see changes in the same order, but the data might not always be the most recent.
- **Eventual Consistency:** Data might not be immediately the same everywhere, but after a while, all copies will become the same. This is faster but may show old data for a short time.

# 5. Consistency Protocols



- Consistency protocols are the rules that help the system decide how and when to update all the copies of data.
- 5.1 Primary-Based Protocols
- One server is the main server (primary), and all updates happen here first.
- After updating, the primary sends the changes to the backup servers (secondaries).
- Example: If you edit a Google Doc, the changes are first saved to the main server, then sent to other servers.



- 5.2 Quorum-Based Protocols
- Instead of updating all copies, the system updates only a majority of them.
- You need to get a certain number of servers (called a quorum) to agree on a change.
- Example: Imagine you need permission from 3 out of 5 people before making a decision.




- 5.3 Gossip Protocols
- Each server talks to a few others, spreading updates like a gossip chain.
- Eventually, all servers will know the latest update, but it may take time.
- Example: Like how a rumor spreads slowly through a school, one person tells another until everyone knows.



## 6. CAP Theorem



- In a distributed system, you can only pick two out of three things at the same time:
- Consistency: All users see the same data at the same time.
- Availability: The system always works, even if some servers are down.
- Partition Tolerance: The system works even if some servers can't communicate with each other.



- You have to sacrifice one of these. For example:
- If you want data to be always available, you might allow temporary inconsistencies.
- If you want strict consistency, your system might go down if some servers can't communicate.

## 7. Real-World Examples

- Google Docs: Uses strong consistency; every change you make is immediately reflected on all devices.
- Amazon DynamoDB: Uses eventual consistency to prioritize fast data access, even if some servers are temporarily outdated.



## 8. Conclusion

- Replica management and consistency protocols help make sure data in a distributed system stays reliable, fast, and available.
- Different systems use different strategies based on whether they prioritize speed, consistency, or fault tolerance.

